

tion-analysis-and-prediction-in-ml

September 9, 2023

```
[2]: # Beer consumption Analysis and prediction using Linear Regression in Machine Learning
import numpy as np
import pandas as pd
beer = pd.read_csv("beer.csv")
```

```
[3]: beer.head(5)
```

```
[3]:      Data  Temperatura Media (C)  Temperatura Minima (C)  \
0  1/1/2015                27,3                23,9
1  1/2/2015                27,02               24,5
2  1/3/2015                24,82               22,4
3  1/4/2015                23,98               21,5
4  1/5/2015                23,82                21

      Temperatura Maxima (C)  Precipitacao (mm)  Final de Semana  \
0                32,5                0                0
1                33,5                0                0
2                29,9                0                1
3                28,6                1,2                1
4                28,3                0                0

      Consumo de cerveja (litros)
0                25.461
1                28.972
2                30.814
3                29.799
4                28.900
```

```
[5]: beer.columns
```

```
[5]: Index(['Data', 'Temperatura Media (C)', 'Temperatura Minima (C)',
        'Temperatura Maxima (C)', 'Precipitacao (mm)', 'Final de Semana',
        'Consumo de cerveja (litros)'],
        dtype='object')
```

dimensions of rows and columns

```
[6]: beer.shape
```

```
[6]: (365, 7)
```

EDA

first, using the normal describe command which just displays the columns of the dataset with the data.

```
[7]: beer.describe
```

```
[7]: <bound method NDFrame.describe of Data Temperatura Media (C)
```

```
Temperatura Minima (C) \
```

0	1/1/2015	27,3	23,9
1	1/2/2015	27,02	24,5
2	1/3/2015	24,82	22,4
3	1/4/2015	23,98	21,5
4	1/5/2015	23,82	21
..
360	12/27/2015	24	21,1
361	12/28/2015	22,64	21,1
362	12/29/2015	21,68	20,3
363	12/30/2015	21,38	19,3
364	12/31/2015	24,76	20,2

	Temperatura Maxima (C)	Precipitacao (mm)	Final de Semana \
0	32,5	0	0
1	33,5	0	0
2	29,9	0	1
3	28,6	1,2	1
4	28,3	0	0
..
360	28,2	13,6	1
361	26,7	0	0
362	24,1	10,3	0
363	22,4	6,3	0
364	29	0	0

	Consumo de cerveja (litros)
0	25.461
1	28.972
2	30.814
3	29.799
4	28.900
..	...
360	32.307
361	26.095
362	22.309

```
363                20.467
364                22.446
```

```
[365 rows x 7 columns]>
```

with describe() command the mathematical description of data inside the dataset is produced namely count, mean of data, its standard deviation, min values, maximum values etc.

```
[8]: beer.describe()
```

```
[8]:          Final de Semana  Consumo de cerveja (litros)
count          365.000000          365.000000
mean           0.284932          25.401367
std            0.452001           4.399143
min            0.000000          14.343000
25%            0.000000          22.008000
50%            0.000000          24.867000
75%            1.000000          28.631000
max            1.000000          37.937000
```

DATA CLEANING

```
[10]: beer.isna()
```

```
[10]:          Data  Temperatura Media (C)  Temperatura Minima (C)  \
0      False                False                False
1      False                False                False
2      False                False                False
3      False                False                False
4      False                False                False
..      ...                ...                ...
360    False                False                False
361    False                False                False
362    False                False                False
363    False                False                False
364    False                False                False

          Temperatura Maxima (C)  Precipitacao (mm)  Final de Semana  \
0                False                False                False
1                False                False                False
2                False                False                False
3                False                False                False
4                False                False                False
..                ...                ...                ...
360              False                False                False
361              False                False                False
362              False                False                False
363              False                False                False
```

```
364                False                False                False
```

```
    Consumo de cerveza (litros)
0                False
1                False
2                False
3                False
4                False
..              ...
360              False
361              False
362              False
363              False
364              False
```

```
[365 rows x 7 columns]
```

```
[11]: # observing if any nan or missing values present or not
beer.isna().any()
```

```
[11]: Data                False
Temperatura Media (C)    False
Temperatura Minima (C)   False
Temperatura Maxima (C)   False
Precipitacao (mm)        False
Final de Semana          False
Consumo de cerveza (litros) False
dtype: bool
```

```
[18]: # dropping any rows present with nan values
beer.dropna(how='all', inplace=True)
```

```
[19]: # replacing commas with period(dot)
beer.replace({' ','.'},regex = True, inplace=True)
```

```
[20]: # converting the type of data to date time
beer['Data'] = pd.to_datetime(beer['Data'])
```

```
[47]: #timestamp = date.map(lambda x : dateutil.parser.parse(x))
#beer = beer.apply(pd.to_numeric)

beer = beer.astype({'Temperatura Media (C)': 'float', 'Temperatura Minima (C)':
    ↪ 'float', 'Temperatura Maxima (C)': 'float', 'Consumo de cerveza (litros)':
    ↪ 'float'})
```

```
[48]: # converting weekdays to weekends
beer.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 365 entries, 0 to 364
Data columns (total 8 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Data                                  365 non-null    datetime64[ns]
1   Temperatura Media (C)                 365 non-null    float64
2   Temperatura Minima (C)                 365 non-null    float64
3   Temperatura Maxima (C)                 365 non-null    float64
4   Precipitacao (mm)                     365 non-null    object
5   Final de Semana                       365 non-null    int64
6   Consumo de cerveja (litros)           365 non-null    float64
7   Day                                   365 non-null    object
dtypes: datetime64[ns](1), float64(4), int64(1), object(2)
memory usage: 22.9+ KB
```

```
[61]: days = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']

beer['Day'] = beer['Day'].replace({'Monday':1, 'Tuesday':2, 'Wednesday':
    ↪3, 'Thursday':4, 'Friday':5, 'Saturday':6, 'Sunday':0})

beer.drop(['Data'], axis=1)
```

```
[61]:
```

	Temperatura Media (C)	Temperatura Minima (C)	Temperatura Maxima (C)	\
0	27.30	23.9	32.5	
1	27.02	24.5	33.5	
2	24.82	22.4	29.9	
3	23.98	21.5	28.6	
4	23.82	21.0	28.3	
..	
360	24.00	21.1	28.2	
361	22.64	21.1	26.7	
362	21.68	20.3	24.1	
363	21.38	19.3	22.4	
364	24.76	20.2	29.0	

	Precipitacao (mm)	Final de Semana	Consumo de cerveja (litros)	Day
0	0	0	25.461	4
1	0	0	28.972	5
2	0	1	30.814	6
3	1.2	1	29.799	0
4	0	0	28.900	1
..
360	13.6	1	32.307	0
361	0	0	26.095	1
362	10.3	0	22.309	2
363	6.3	0	20.467	3

364	0	0	22.446	4
-----	---	---	--------	---

```
[365 rows x 7 columns]
```

```
[64]: #beer['Day']=beer['Data'].apply(lambda a: days[a.weekday()])
beer.head(5)
```

[64]:	Data	Temperatura Media (C)	Temperatura Minima (C)	\
0	14200704000000000000	27.30	23.9	
1	14201568000000000000	27.02	24.5	
2	14202432000000000000	24.82	22.4	
3	14203296000000000000	23.98	21.5	
4	14204160000000000000	23.82	21.0	

	Temperatura Maxima (C)	Precipitacao (mm)	Final de Semana \
0	32.5	0	0
1	33.5	0	0
2	29.9	0	1
3	28.6	1.2	1
4	28.3	0	0


	Consumo de cerveza (litros)	Day
0	25.461	4
1	28.972	5
2	30.814	6
3	29.799	0
4	28.900	1

```
[65]: beer.columns
```

```
[65]: Index(['Data', 'Temperatura Media (C)', 'Temperatura Minima (C)',
          'Temperatura Maxima (C)', 'Precipitacao (mm)', 'Final de Semana',
          'Consumo de cerveja (litros)', 'Day'],
          dtype='object')
```

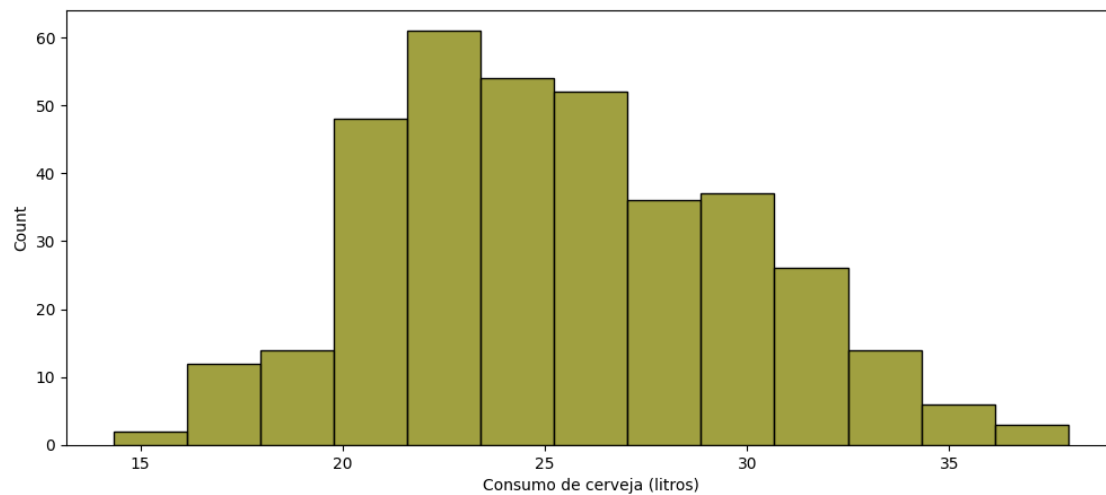
```
[66]: import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(11,6))
ax = sns.boxplot(data=beer, x='Day', y='Consumo de cerveja (litros)',
                color="indianred")
```




```
[112]: plt.figure(figsize=(12,5))
sns.histplot(beer , x="Consumo de cerveja (litros)" , color='olive')
```

```
[112]: <Axes: xlabel='Consumo de cerveja (litros)', ylabel='Count'>
```



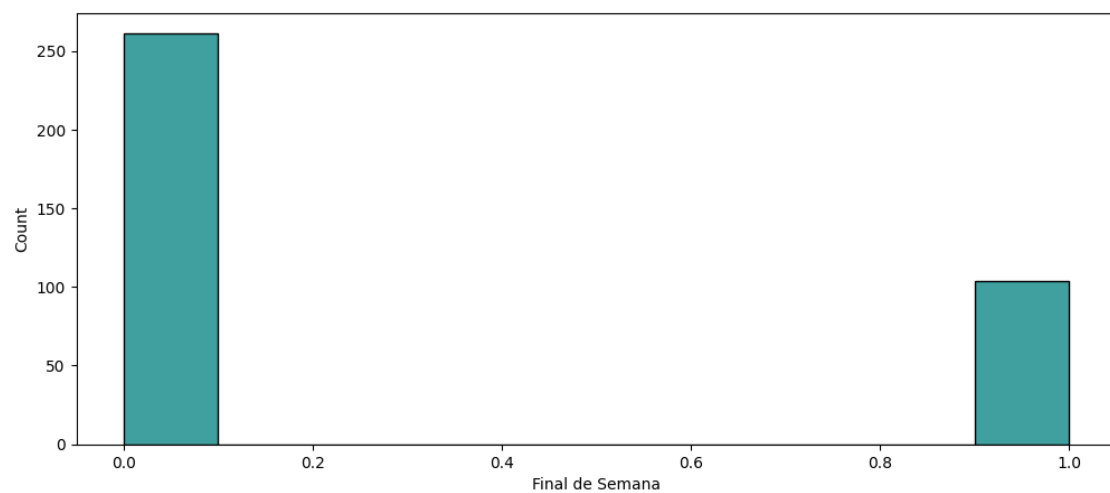
```
[127]: plt.figure(figsize=(23,5))
sns.histplot(beer , x="Precipitacao (mm)" , color='violet')
```

```
[127]: <Axes: xlabel='Precipitacao (mm)', ylabel='Count'>
```




```
[117]: plt.figure(figsize=(12,5))
sns.histplot(beer , x="Final de Semana" , color='teal')
```

```
[117]: <Axes: xlabel='Final de Semana', ylabel='Count'>
```




```
[126]: plt.figure(figsize=(29,5))
sns.lineplot(x='Precipitacao (mm)', y='Day', data=beer, color='limegreen')
```

```
[126]: <Axes: xlabel='Precipitacao (mm)', ylabel='Day'>
```




```
[129]: plt.figure(figsize=(19,5))
sns.lineplot(x='Precipitacao (mm)', y='Final de Semana', data=beer,
            color='saddlebrown')
```


```
[129]: <Axes: xlabel='Precipitacao (mm)', ylabel='Final de Semana'>
```



```
[133]: plt.figure(figsize=(10,5))
sns.lmplot(x='Consumo de cerveja (litros)', y='Final de Semana', data=beer)
```

```
[133]: <seaborn.axisgrid.FacetGrid at 0x79a78284a9e0>
```

```
<Figure size 1000x500 with 0 Axes>
```



```
[85]: from sklearn.model_selection import train_test_split

# dropping the target values
x = beer.drop("Data", axis=1)

y= beer['Consumo de cerveza (litros)']

xtrain, xtest, ytrain, ytest = train_test_split(x,y, test_size=0.2,
↳random_state=42)
```

```
[86]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import confusion_matrix
from sklearn import metrics

L=LinearRegression()
beer['Data'] = pd.to_numeric(pd.to_datetime(beer['Data']))
L.fit(xtrain,ytrain)
```

```
[86]: LinearRegression()
```

```
[87]: y_pred = L.predict(xtest)
      print(L.score(xtest , ytest))
```

1.0

```
[102]: from sklearn.model_selection import train_test_split
      x = beer.drop("Consumo de cerveza (litros)", axis=1)
      y= beer['Consumo de cerveza (litros)']

      xtrain, xtest, ytrain, ytest = train_test_split(x,y, test_size=0.2,
      ↪random_state=32)

      from sklearn.linear_model import LinearRegression
      from sklearn.metrics import confusion_matrix
      from sklearn import metrics

      L=LinearRegression()
      beer['Data'] = pd.to_numeric(pd.to_datetime(beer['Data']))
      L.fit(xtrain,ytrain)
      y_pred = L.predict(xtest)
      print(L.score(xtest , ytest))
```

0.44753330684957426

```
[70]: y_pred
```

```
[70]: array([26.36476346, 27.38357735, 29.51945925, 24.24788941, 28.41127772,
        20.49956302, 25.06477236, 21.15087143, 20.76139266, 22.75227083,
        28.39966701, 27.80594359, 22.59541924, 26.80878518, 22.9637712 ,
        24.86539129, 30.0535599 , 23.1264367 , 26.60473775, 24.73355932,
        30.1660252 , 22.39744558, 26.00566154, 20.38451303, 21.48328967,
        25.64000457, 24.74727957, 29.82904921, 25.92465065, 26.8759044 ,
        25.33346274, 26.53546932, 30.11199202, 28.36456217, 24.78968375,
        23.91899357, 28.33134898, 25.17047699, 27.31096035, 27.70890767,
        28.18383108, 28.46700732, 24.95220103, 22.04914598, 26.11424429,
        25.78056654, 26.43378333, 25.5369827 , 27.02452207, 25.53011284,
        24.13615029, 27.63571748, 27.54326356, 28.81473754, 21.21265299,
        27.30339868, 26.79924668, 27.38626656, 21.47241743, 25.90397454,
        23.69932148, 29.55983883, 25.91833692, 24.55715151, 29.53676997,
        26.01200698, 25.32848714, 22.0456233 , 21.87785279, 25.30048265,
        26.29560742, 24.79553862, 26.67603486])
```


```
[72]: ytest
```

```
[72]: 193    25.640
      33    31.655
```

```
15      29.938
309     20.575
57      26.241
...
203     19.029
82      21.617
94      32.713
192     33.298
325     31.933
```

Name: Consumo de cerveza (litros), Length: 73, dtype: float64

```
[77]: sns.lmplot(x="Day", y="Consumo de cerveza (litros)", data = beer, order = 2)
plt.show()
```



```
[101]: regr = LinearRegression()
regr.fit(xtrain, ytrain)
print(regr.score(xtest, ytest))
y_pred = regr.predict(xtest)
```

0.531384628597559

```
[103]: from sklearn.metrics import mean_absolute_error, mean_squared_error

mae = mean_absolute_error(y_true=ytest, y_pred=y_pred)

#squared True returns MSE value, False returns RMSE value.
mse = mean_squared_error(y_true=ytest, y_pred=y_pred) #default=True
rmse = mean_squared_error(y_true=ytest, y_pred=y_pred, squared=False)

print("MAE:", mae)
print("MSE:", mse)
print("RMSE:", rmse)
```

MAE: 2.6345518760119706
MSE: 10.359020220783298
RMSE: 3.2185431829918483

```
[110]: from sklearn.metrics import confusion_matrix #confusion matrix only
        ↳ displays info of numeric data
correlation_metrics = beer.corr()
fig = plt.figure(figsize = (8,8))
sns.heatmap(correlation_metrics , vmax = 0.9 , square = True, annot=True)
plt.show()
```

<ipython-input-110-4dbc9bc8b4a6>:2: FutureWarning: The default value of numeric_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric_only to silence this warning.

```
correlation_metrics = beer.corr()
```

